

## GRADIENT DESCENT

### \* Sample, Batch, Epoch

- A sample is a single row of data.
- The batch size defines the number of samples to work through, before updating the internal model parameters.
  - Depending on the batch size different learning algos are defined.
- The number of epochs is a hyperparameter that defines the number of times the learning algo will work through the entire training dataset.
  - One epoch

### \* An epoch

one epoch:

1. Randomly divide training set into  $m = N/k$  batches.
2. Use a batch of training samples to compute  $J(w)$ .
3. Update  $w$  as:  $w^{t+1} = w^t - \eta \frac{\partial J}{\partial w}$
4. Repeat 2 & 3 using different subsets all samples are used once.

Q. what is the size of the batch?

- $1, \dots, k, \dots, N$
- May depend on hardware.
- we repeat epochs until convergence.

$1, \dots, k$

$k+1, \dots, 2k$

$2k+1$

## \* Batch Gradient Descent

→ Training set: Ex: ImageNet has 14M images

→ Approach:

- Compute the loss  $J(w)$  on the entire training, update the parameters  $w$ .
- At the next epoch, shuffle the training data, and repeat above process.

→ Typical batch size: Size of the training set.

## \* Mini-batch Gradient Descent

→ It is wasteful to compute the loss over the entire set to perform a single parameter update for large datasets.

- Ex: ImageNet has 14M images

- GD is replaced with mini-batch GD

→ Mini-batch ~~GD~~ GD


• Approach:

- Compute the loss  $L(w)$  on a batch of images, update the parameters  $w$

How to learn  $w$  &  $b$ ? How will our model learn?

learn loss fn  
gradient descent

How many neurons  
in one layer.



$$s = \sum_{i=1}^n x_i w_i + b \quad y = \phi(s)$$

loss / cost / objective function

- supervised learning.
- The loss function provides the cost of being wrong, by measuring the quality of a particular set of parameters based on how well the output of the network agrees with the ground ~~label~~ both labels in the training data.

$L(\theta) = \text{distance} \left( \overset{\substack{\text{Input, features} \\ \text{error}}}{f_{\theta}(x)}, \overset{\substack{\text{label, ground truth} \\ \text{label (true)}}}{y} \right)$   
 difference b/w actual & predicted  
 $\theta = \text{parameters (weights, biases)}$   
 $y_{\text{pred}} = f_{\theta}(x)$  learned by the model

learning process



$$y = w_0 + w_1 x_1 + w_2 x_2$$

depending upon  
feedback.

start random value (for slope).

1. Start with random values of  $w_i$  (training data)
2. Evaluate the goodness of the line, determined with a loss function,  $J(w)$ .
3. The weights  $w_i$  are changed acc moving the line to a better position.  
→ Note:  $J(w)$  should be min when the training samples are correctly classified.
4. Repeat 2 & 3 until  $J(w) < \gamma$

Gradient Descent → linear reg<sup>n</sup> / NN / CNN / RN

learns parameters.

low loss function: Gradient Descent in Action.



Force you to go to dir<sup>n</sup>  
in which max descent is  
less.